

Languages: Python, TypeScript, JavaScript, Java, SQL, C/C++
AI Systems: Agentic AI, agent orchestration & automation, MCP servers/tooling, prompt caching, RAG, embedding generation, vector databases, Elasticsearch, LiteLLM, Vercel AI SDK, Langfuse, LlamaIndex
Backend & Infra: API protocols (REST, GraphQL, RPC), Node.js, Express.js, SQL, data modeling, ORMs, OAuth, webhooks, async event orchestration (Hatchet), message queues (RabbitMQ), caching/streaming (Redis)
Frontend: React, Redux, data visualization, Retool dashboards
Security & Enterprise: ACL/RBAC/IAM, entitlement & access control, credential isolation, JWT, multi-tenant authorization

PROFESSIONAL EXPERIENCE

Kindo AI • Enterprise AI Automation Platform
Senior Software Engineer May 2024 – Present

- Delivered a production multi-agent orchestration system, enabling hierarchical agent spawning, scoped context injection, sub-agent output guardrails, and permission-aware execution for enterprise workflows.
- Launched a distributed AI orchestration system (Hatchet framework), moving LLM request/response lifecycles from synchronous request flows to durable, retryable, long-running workflows; contributed to cross-functional Agent Runtime initiative rearchitecting the platform's stateful execution substrate, encompassing persistent workspaces, agent skill registries, memory compaction, and multi-tenant sandbox infrastructure.
- Led 0-to-1 platform initiatives, including Agent Version Control and asynchronous Scheduled Agents, shipping version snapshots, restore flows, run-level version stamping, audit logging, cron-based automation, validation, and org-level guardrails from backend through UI.
- Designed Kindo's Native Integrations Provider, replacing a third-party platform and consolidating OAuth, webhooks, credential management, provider routing, and MCP integration workflows.
- Expanded secure multi-tenant access controls across agent execution and platform features, including user group-scoped credentials, admin-gated chat/file/API-key entitlements, refined integration sharing permissions, shared-agent tool availability warnings, and MCP user-context propagation.
- Built Command Center analytics for agent and credit usage, with zero-downtime migrations, aggregation APIs, model attribution, workflow-run breakdowns, and drill-down time-series dashboards.

Roses Del Fuego • Luxury Flower Startup
Freelance Web Developer Jan. 2024 – May 2024

- Developed a custom Shopify administrative application with complex queries, advanced scheduling features, and custom data fields to support supply-chain workflows for a luxury floral business.
- Built multilingual, stylistically accessible UX/UI using React, Shopify API, GraphQL, JavaScript, Liquid, SQLite.

Ford Motor Company • Automotive Manufacturing Leader
Product Development Engineering Intern May 2022 – Aug. 2022

- Developed a machine learning model to select representative high-frequency Connected Vehicle data cycles and predict customer vehicle behavior and EV charging patterns.
- Wrangled data from 50,000+ unique vehicle identification numbers, prioritizing efficient preprocessing, extensible algorithm design, and scalable model implementation.

EDUCATION

Columbia University in the City of New York
M.S. in Computer Science, Machine Learning • GPA: 3.8 • Lustgarten-Whitney Family Fellowship Dec. 2023
University of California, Los Angeles
B.S. in Cognitive Science, Specialization in Computing • GPA: 3.95 June 2021
B.A. in Human Biology & Society • GPA: 3.95 June 2021